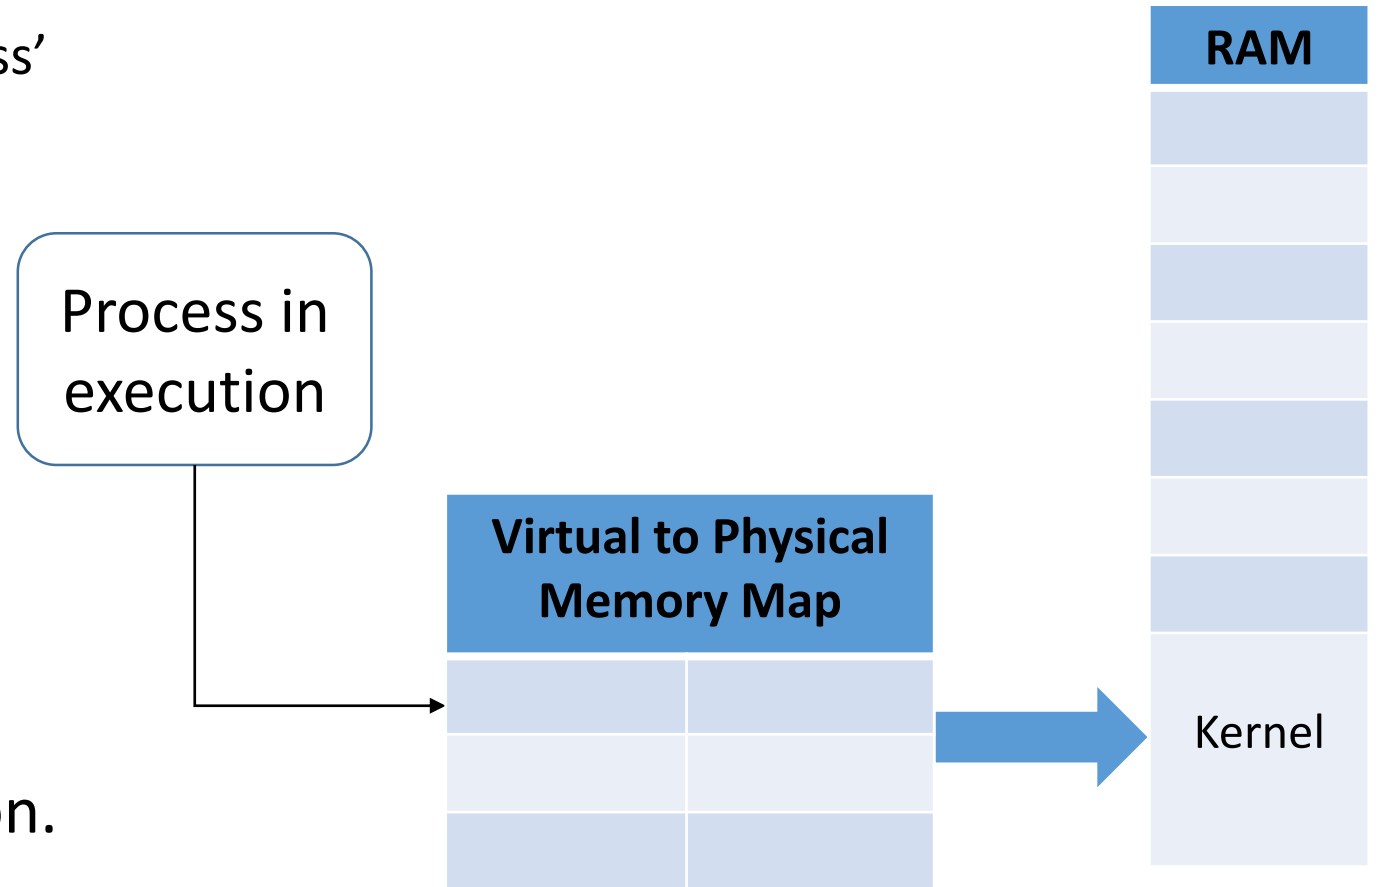


Inter Process Communication (IPC) Process Synchronization Deadlocks

CS3003D: Operating Systems

Need for IPC

- Each process has it's own virtual address space
 - Cannot view/access another process' address space
 - MMU maps the virtual address to Physical/RAM address
- Cannot guess/determine the physical address mapping
- How does one process communicate with another process?
Inter Process Communication.


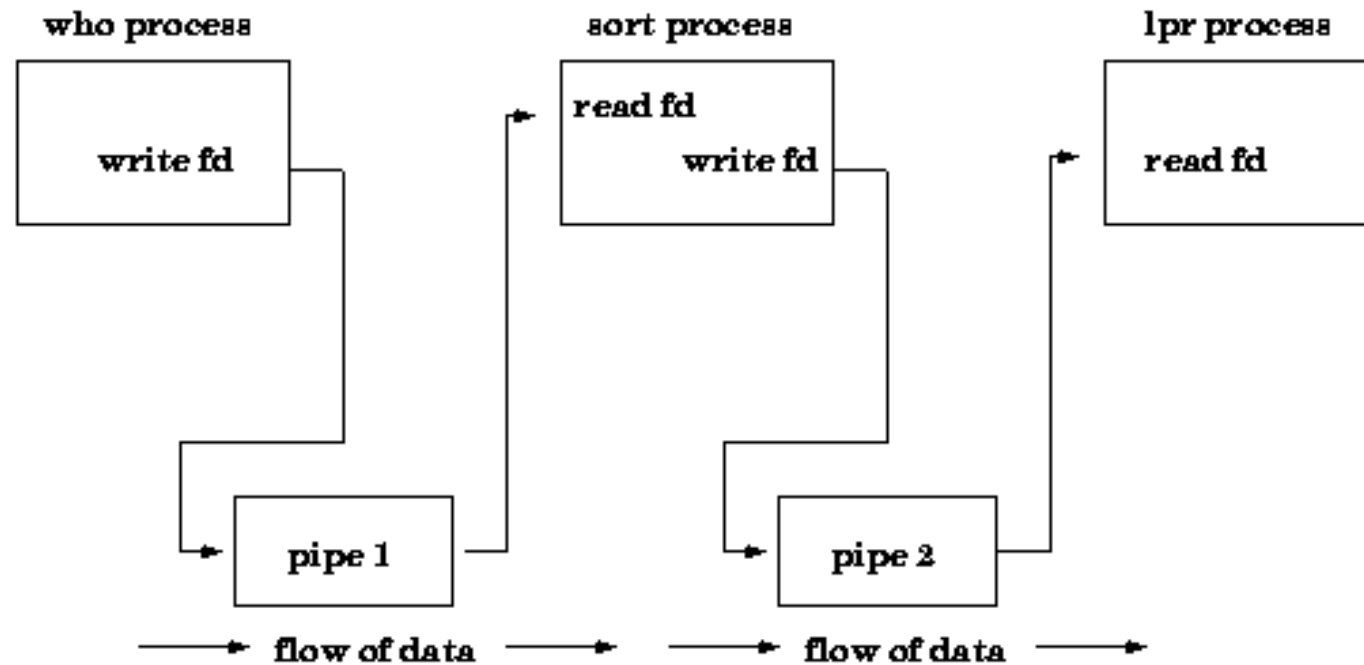


IPC

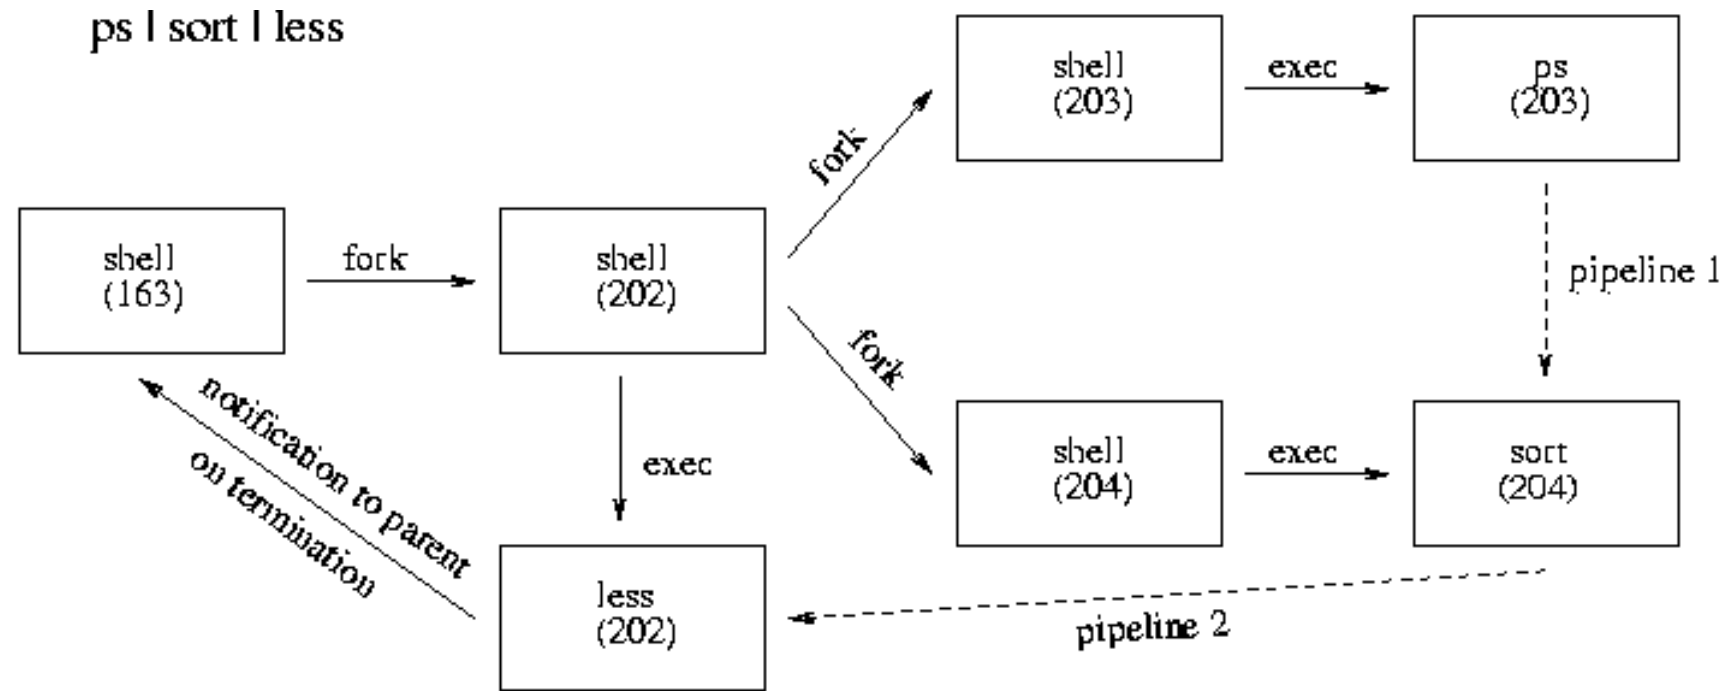
- Information sharing between processes
 - P1: data collection from the environment
 - P2: analyses the data collected
 - P3: actuates the external devices
- Convenient usage

Pipes (unnamed pipes)

- `who | sort | lpr`





ps | sort | less



Pipes

- Always communication between child and parent processes
- Unidirectional
- `fd[0]` – reads from the pipe
- `fd[1]` – writes to a pipe

- Messages from parent to child
 - parent closes `fd[0]`
 - child closes `fd[1]`
- Messages from child to parent
 - parent closes `fd[1]`
 - child closes `fd[0]`



A pipe program

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main(void)
{
    int      fd[2], nbytes;
    pid_t    childpid;
    char      string[] = "Hello, world!\n";
    char      readbuffer[80];

    pipe(fd);

    if((childpid = fork()) == -1)
    {
        perror("fork");
        exit(1);
    }
```

```
    if(childpid == 0)
    {
        /* Child process closes up input side of pipe */
        close(fd[0]);

        /* Send "string" through the output side of pipe */
        write(fd[1], string, (strlen(string)+1));
        exit(0);
    }
    else
    {
        /* Parent process closes up output side of pipe */
        close(fd[1]);

        /* Read in a string from the pipe */
        nbytes = read(fd[0], readbuffer, sizeof(readbuffer));
        printf("Received string: %s", readbuffer);
    }

    return(0);
}
```

Named Pipe

- Also known as FIFO
- A named pipe can last as long as the system is up, beyond the life of the process
 - It can be deleted if no longer used
- Usually named pipe appears as a file; two or more processes can communicate by reading/writing from/to the file
- The named pipe resides in the kernel and not on a physical file system
- Function call:

```
int mkfifo(const char *pathname, mode_t mode);  
mknod()
```


Named Pipe ... Contd.

Process A

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <unistd.h>

#include <linux/stat.h>

#define FIFO_FILE      "MYFIFO"

int main(void)
{
    FILE *fp;
    char readbuf[80];

    /* Create the FIFO if it does not exist */
    umask(0);
    mknod(FIFO_FILE, S_IFIFO|0666, 0);

    while(1)
    {
        fp = fopen(FIFO_FILE, "r");
        fgets(readbuf, 80, fp);
        printf("Received string: %s\n", readbuf);
        fclose(fp);
    }

    return(0);
}
```

Process B

```
#include <stdio.h>
#include <stdlib.h>

#define FIFO_FILE      "MYFIFO"

int main(int argc, char *argv[])
{
    FILE *fp;

    if ( argc != 2 ) {
        printf("USAGE: fifoclient [string]\n");
        exit(1);
    }


    if((fp = fopen(FIFO_FILE, "w")) == NULL) {
        perror("fopen");
        exit(1);
    }

    fputs(argv[1], fp);

    fclose(fp);
    return(0);
}
```

Shared Memory

- Process creates an area in RAM so that other processes can access
- Reading from and writing to the shared memory space
- There is no intermediary such as message queues, pipes, etc.
- Adv.: Fast
- Prone to error; processes should synchronize



Shared Memory ... Contd.

- `shmget()` is used to obtain access to a shared memory segment.
- `int shmget(key_t key, size_t size, int shmflg);`
 - `key` : unique key/ID
 - `size` : size in bytes of the requested shared memory
 - `shmflg` : initial access permissions and creation control flags
 - returns ID of the segment: `shmID`
- `shmat()` and `shmdt()` are used to attach and detach shared memory segments
 - Attach a process (address space of the process) to the shared memory, `shmID`
 - Detach shared memory

Shared Memory – Process01

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
#include <stdio.h>

#define SHMSZ    27 // size of shared memory

main()
{
    char c;
    int shmid;
    key_t key;
    char *shm, *s;

    /* We'll name our shared memory segment "5678". */
    key = 5678;

    /* Create the segment. */
    if ((shmid = shmget(key, SHMSZ, IPC_CREAT | 0666)) < 0) {
        perror("shmget");
        exit(1);
    }

    /* Now attach the segment to the data space. */
    if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {
        perror("shmat");
        exit(1);
    }

    /* Now put some things into the memory for the other process to read. */
    s = shm;
    for (c = 'a'; c <= 'z'; c++)
        *s++ = c;
    *s = NULL;

    /*
     * Finally, we wait until the other process
     * changes the first character of our memory to '*',
     * indicating that it has read the memory area */
    while (*shm != '*')
        sleep(1);

    exit(0);
}
```

Source: <https://users.cs.cf.ac.uk/Dave.Marshall/C/node27.html>

Shared Memory – Process02

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
#include <stdio.h>

#define SHMSZ    27

main()
{
    int shmid;
    key_t key;
    char *shm, *s;

    /* get the segment named "5678", created by the server. */
    key = 5678;

    /* Locate the segment. */
    if ((shmid = shmget(key, SHMSZ, 0666)) < 0) {
        perror("shmget");
        exit(1);
    }

    /* Attach the segment to the data space. */
    if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {
        perror("shmat");
        exit(1);
    }
}
```

```
/* Now read what the server put in the memory */
for (s = shm; *s != NULL; s++)
    putchar(*s);
putchar('\n');


/*
 * Finally, change the first character of the
 * segment to '*', indicating we have read
 * the segment.
 */
*shm = '*';

exit(0);
}
```

Source: <https://users.cs.cf.ac.uk/Dave.Marshall/C/node27.html>

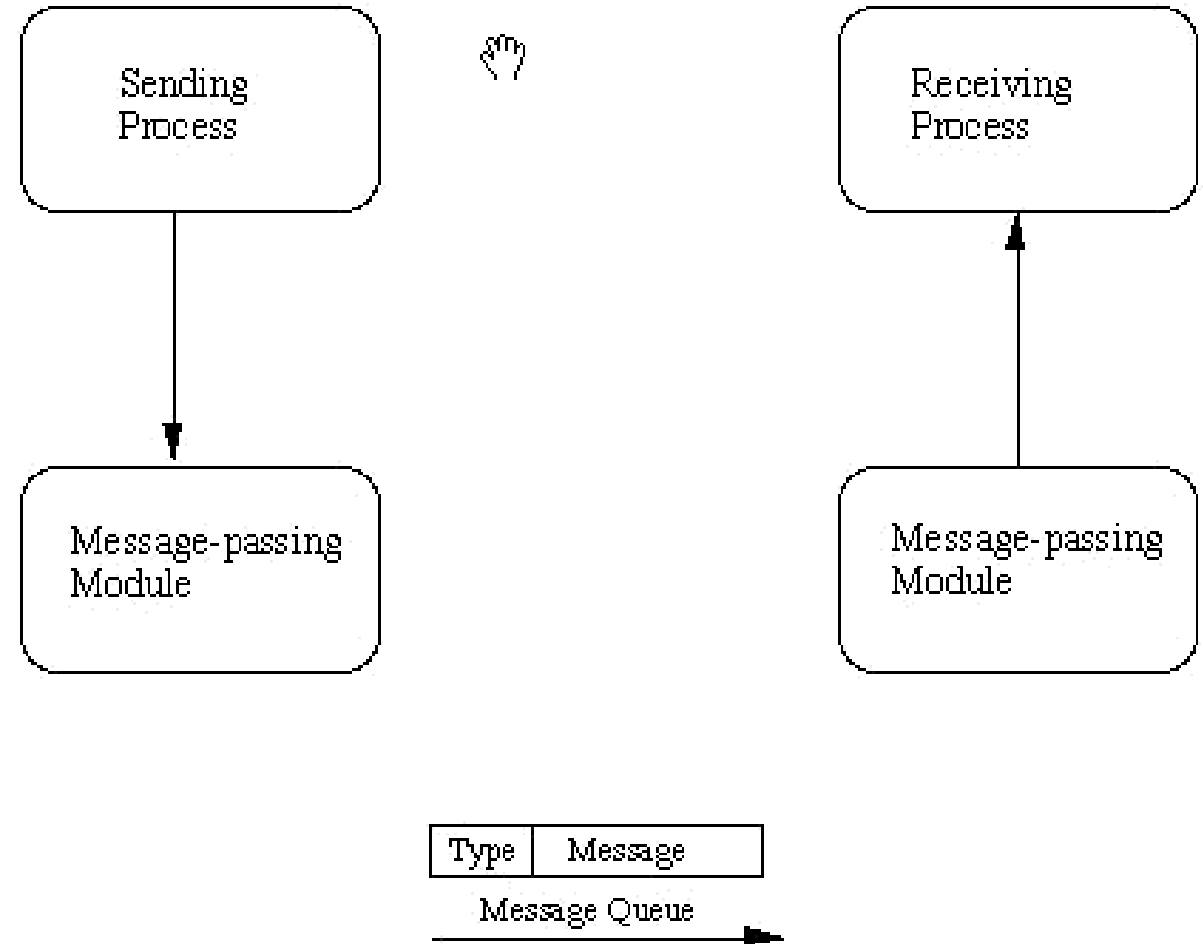
Message Passing

- Shared memory – reading and writing
- Message passing – send and receive
- System calls `send()` and `receive()` are used for communication
- Less error prone
- Slow because of system calls



Message Queues

- Easy implementation when communication between two or more processes



- `int msgget(key_t key, int msgflg)`
creates or allocates a message queue
- `int msgsnd(int msgid, const void *msgp, size_t msgsz, int msgflg)`
system call sends/appends a message into the message queue
- `int msgrcv(int msgid, const void *msgp, size_t msgsz, long msgtype, int msgflg)`
system call retrieves the message from the message queue

Ref.: <https://users.cs.cf.ac.uk/Dave.Marshall/C/node25.html>


```

// C Program for Message Queue (Writer Process)
#include <stdio.h>
#include <sys/ipc.h>
#include <sys/msg.h>

// structure for message queue
struct mesg_buffer {
    long mesg_type;
    char mesg_text[100];
} message;

int main()
{
    key_t key;
    int msgid;

    // ftok to generate unique key
    key = ftok("progfile", 65);

    // msgget creates a message queue and returns identifier
    msgid = msgget(key, 0666 | IPC_CREAT);
    message.mesg_type = 1;

    printf("Write Data: ");
    gets(message.mesg_text);

    // msgsnd to send message
    msgsnd(msgid, &message, sizeof(message), 0);

    // display the message
    printf("Data send is : %s \n", message.mesg_text);

    return 0;
}

```

```

// C Program for Message Queue (Reader Process)
#include <stdio.h>
#include <sys/ipc.h>
#include <sys/msg.h>

// structure for message queue
struct mesg_buffer {
    long mesg_type;
    char mesg_text[100];
} message;

int main()
{
    key_t key;
    int msgid;

    // ftok to generate unique key
    key = ftok("progfile", 65);

    // msgget creates a message queue and returns identifier
    msgid = msgget(key, 0666 | IPC_CREAT);

    // msgrcv to receive message
    msgrcv(msgid, &message, sizeof(message), 1, 0);

    // display the message
    printf("Data Received is : %s \n",
        message.mesg_text);

    // to destroy the message queue
    msgctl(msgid, IPC_RMID, NULL);

    return 0;
}

```

Synchronization

shared variable
int flag = 5

ProgramA

flag++

Output value of flag can be 5, 4, or 6 based on the way the processes are executing, when context switching happens

ProgramB

flag --

- 1) reg1 = flag
- 2) reg1 = reg1 + 1
- 3) flag = reg1

- 4) reg2 = flag
- 5) reg2 = reg2 - 1
- 6) flag = reg2

Scenario1

ProcessA

- 1)
- 2)
- 3) flag = 6

Context Switch

ProcessB

- 4)
- 5)
- 6) flag = 5

Scenario2

ProcessB

- 1)
- 2)
- 3) flag = 4

Context Switch

ProcessA

- 4)
- 5)
- 6) flag = 5

Scenario3

ProcessA

- 1) reg1=5

Context Switch

ProcessB

- 2) reg2 = 5
- 3) reg2 = 4
- 4) flag = 4

Context Switch

ProcessA

- 5) reg1 = 6
- 6) flag = 6

Scenario4

ProcessB

- 1) reg2=5

Context Switch

ProcessA

- 2) reg2 = 5
- 3) reg2 = 6
- 4) flag = 6

Context Switch

ProcessB

- 5) reg1 = 4
- 6) flag = 4

Race condition

- Many processes manipulate the same data portion
- During concurrent execution outcome depends upon the order in which the access happens
- Incorrect data leads to misleading output
- Can be prevented by synchronization between processes

How to avoid race condition?

- Prohibit more than one process from reading and writing the shared data (critical section) at the same time

Three requirements for critical section problem

- Mutual Exclusion

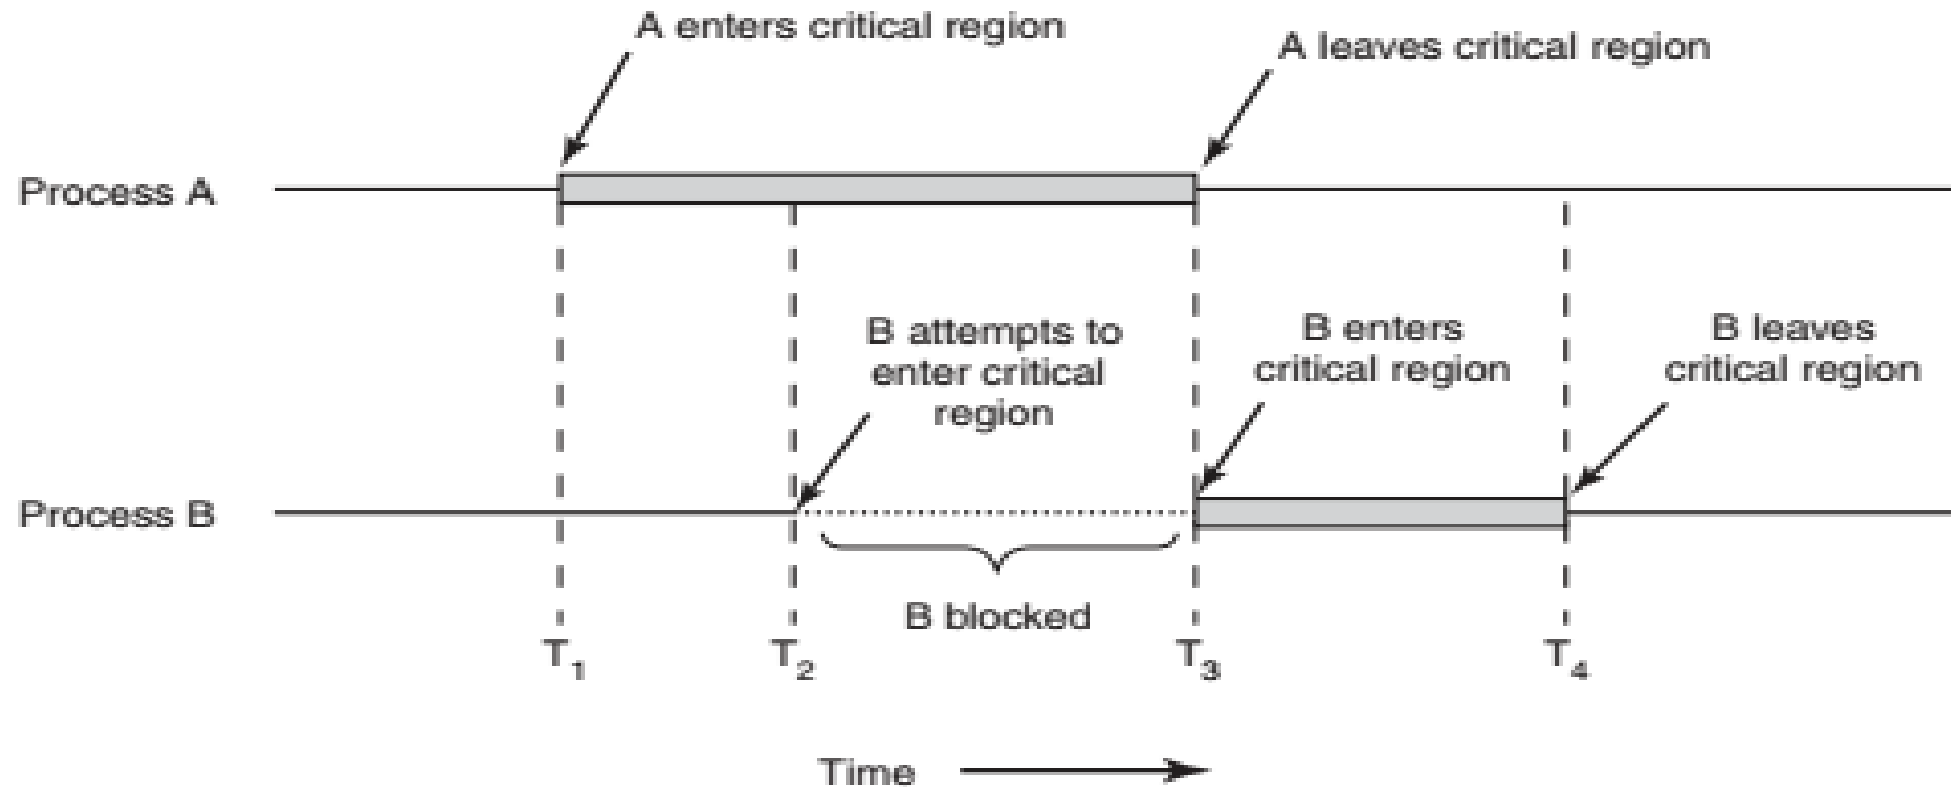
- No two processes may be simultaneously inside their critical regions

- Progress

- No process running outside its critical region may block any process

- No starvation (bounded waiting)

- No process should have to wait forever to enter its critical region



Solutions to critical section

- Disable interrupts

- Context switches will not happen
- Codes that execute in the kernel can only disable interrupts
- User processes/application programs cannot disable interrupts

```
While(TRUE) {  
    // code area  
    disable_interrupts() < LOCK  
    critical_section  
    enable_interrupts() < UNLOCK  
    // other code area  
}
```

Busy waiting

Process-1

```
while (TRUE) {  
    while (turn == 2);    // LOCK  
    critical_section  
    turn = 2;            // UNLOCK  
}
```

Shared
int turn = 1;

Process-2

```
while (TRUE) {  
    while (turn == 1);    // LOCK  
    critical_section  
    turn = 1;            // UNLOCK  
}
```

- Mutual exclusion achieved
- Busy waiting – resource wastage
 - When Process-2 executes first, always in loop; always in primary memory – either at READY state or at RUNNING state
- Progress condition is violated
 - Process-1 -> Process-2 -> Process-1 -> Process-2 ->

No Mutual Exclusion

Shared

p1_inside = false, p2_inside = false

Process-1

```
while (TRUE) {  
    while (p2_inside == TRUE);    // LOCK  
    p1_inside = TRUE;  
    critical_section  
    p1_inside = FALSE;           // UNLOCK  
}
```

Process-2

```
while (TRUE) {  
    while (p1_inside == TRUE);    // LOCK  
    p2_inside = TRUE;  
    critical_section  
    p2_inside = FALSE;           // UNLOCK  
}
```

```
while(p2_inside == TRUE);  
// Context Switch (Process-2)  
while(p1_inside == TRUE);  
p2_inside = TRUE;  
// Context Switch (Process-1)  
p1_inside = TRUE;
```

- Mutual exclusion is not guaranteed
- Both processes can enter into critical section

Endless Wait

Shared

p1_inside = false, p2_inside = false

Process-1

Process-2

```
while (TRUE) {  
    p1_inside = TRUE;  
    while (p2_inside == TRUE);    // LOCK  
    critical_section  
    p1_inside = FALSE;           // UNLOCK  
}
```

```
while (TRUE) {  
    p2_inside = TRUE;  
    while (p1_inside == TRUE);    // LOCK  
    critical_section  
    p2_inside = FALSE;           // UNLOCK  
}
```

```
p1_inside = TRUE  
// Context Switch (Process-2)  
p2_inside = TRUE;
```

- Achieves Mutual exclusion
- Can it progress?
 - DEADLOCK!

Peterson's Solution

Process-1

Shared

p1_inside, p2_inside, favoured

Process-2

```
while (TRUE) {  
    p1_inside = TRUE;  
    favoured = 2;  
    while (p2_inside == TRUE AND favoured =2); // LOCK  
    critical_section  
    p1_inside = FALSE;           // UNLOCK  
}
```

```
while (TRUE) {  
    p2_inside = TRUE;  
    favoured = 1;  
    while (p1_inside == TRUE AND favoured =1); // LOCK  
    critical_section  
    p2_inside = FALSE;           // UNLOCK  
}
```

Breaking the deadlock as one of the processes is favoured

Bakery Algorithm – synchronization between N processes ($N > 2$)

- Proposed by Leslie Lamport
- Similar to token system in the bakeries/banks
 - Customers upon entering the bank is issued with the token
 - Waits until his/her turn arrives
 - Dispense the token and the service is rendered

Ref.: <http://www.cs.umd.edu/~shankar/412-S99/note-7.html>

Simplified version of Bakery algorithm

- Each process is numbered 0 to N-1
- Each process i has an integer variable num[i], initially 0, that is readable by all processes but writeable by process i only

Entry(i) {

 num[i] = MAX(num[0], num[1], ... , num[N-1]) + 1 ;

 for p = 0 to N-1 do {

 while (num[p] != 0 AND num[p] < num[i]) do no-op ;

 }

}

Lock

Critical section

Exit(i) {

 num[i] = 0 ;

}

Unlock

Example

num[i]	P0	P1	P2	P3	P4
Initial	0	0	0	0	0
P2	0	0	1	0	0
P3	0	0	0	2	0
P4	0	0	0	0	3
P0	4	0	0	0	0
P1	4	5	0	0	0
Final	4	5	1	2	3

$$\text{num}[i] = \text{MAX}(\text{num}[0], \text{num}[1], \dots, \text{num}[N-1]) + 1$$

num[i]	P0	P1	P2	P3	P4
Initial	4	5	1	2	3
P2	4	5	0	2	3
P3	4	5	0	0	3
P4	4	5	0	0	0
P0	0	5	0	0	0
P1	0	0	0	0	0
Final	0	0	0	0	0

```
for p = 0 to N-1 do {  
    while(num[p] != 0 AND num[p] < num[i]) do no-op ;  
}
```

Problem!

Assumption: No two processes get the same token

When two process gets the same num[i] value (same token)

Two processes enter into the critical section

Bakery algorithm

MAX() is no more assumed to be atomic

Introduction of an array of N booleans, choosing[i] = FALSE

```
Entry(i) {  
    choosing[i] = TRUE;  
    num[i] = MAX( num[0], num[1], ... , num[N-1] ) + 1 ;  
    choosing[i] = FALSE;  
    for p = 0 to N-1 do {  
        while(choosing[p]);  
        while (num[p] != 0 AND (num[p],p) < (num[i],i)) do no-op ;  
    }  
}
```

(a,b) < (c,d) is equivalent to (a<c) or ((a==c) AND (b<d))

Lock

Critical section

```
Exit(i) {  
    num[i] = 0 ;  
}
```

Unlock

Example

num[i]	P0	P1	P2	P3	P4
Initial	0	0	0	0	0
P2	0	0	1	0	0
P3	0	0	0	2	0
P4	0	0	0	0	2
P0	3	0	0	0	0
P1	3	4	0	0	0
Final	3	4	1	2	2

num[i]	P0	P1	P2	P3	P4
Initial	3	4	1	2	2
P2	3	4	0	2	2
P3	3	4	0	0	2
P4	3	4	0	0	0
P0	0	4	0	0	0
P1	0	0	0	0	0
Final	0	0	0	0	0

$\text{num}[i] = \text{MAX}(\text{num}[0], \text{num}[1], \dots, \text{num}[N-1]) + 1$

```
for p = 0 to N-1 do {
    while(num[p] != 0 AND num[p] < num[i]) do no-op ;
}
```

In the case of a tie in num[i],
 $(\text{num}[p], p) < (\text{num}[i], i)$ ensures that the process
 with lesser ID prevails

Hardware Locks

```
lock = 0

Process-1
while(1) {
    while(lock != 0);
    lock = 1;    // lock
    // critical_section
    lock = 0;    // unlock
}

Process-2
while(1) {
    while(lock != 0);
    lock = 1;    // lock
    // critical_section
    lock = 0;    // unlock
}
```

```
lock = 0
Process-1
while(lock != 0)
    Context Switch
Process-2
while(lock != 0)
    lock = 1
    Context Switch
Process-1
lock = 1
Both processes in critical section
```

Mutual exclusion not possible

```
while(lock != 0);    /* should be made
lock = 1;            atomic */
```

Test and Set

```
int test_and_set(int *L) {  
    int prev = *L;  
    *L = 1;  
    return prev;  
}
```

```
while(1) {  
    while (test_and_set(&lock) == 1);  
    // critical_section  
    lock = 0;  
}
```



- An atomic function
 - Only one process can access the test_and_set function
 - Hardware ensures only one process can execute the TAS at a time before another process wants to execute
-
- test_and_set will read lock=0 and set lock=1
 - Other processes will see lock=0 and infinitely be in the while loop.

xchg instruction ... Intel equivalent for test and set

```
int xchg(int *L, int val) {  
    int prev = *L;  
    *L = val;  
    return prev;  
}
```

- It is an atomic instruction

```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}
```

```
void acquire(int *locked) {  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
    }  
}
```

```
void release(*locked) {  
    *locked = 0;  
}
```

Spinlock

Process1

Acquire(&locked)
Critical_section
Release(&locked)

Process2

Acquire(&locked)
Critical_section
Release(&locked)

- Only one process can acquire the lock
- Meanwhile other processes wait in a loop for the lock
- When a process releases the lock, it becomes available for other processes

Mutex

- Spinlock is good at short critical sections such as counter increment, accessing array element, etc.
- When period of wait is longer, spinlock is not that effective
- Alternative: Mutex

Mutex

- sleep() – running state to block state
- wakeup() – block state to ready queue

```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}
```

```
void lock(int *locked) {  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
        else  
            sleep()  
    }  
}
```

```
void unlock(*locked) {  
    locked = 0;  
    wakeup();  
}
```

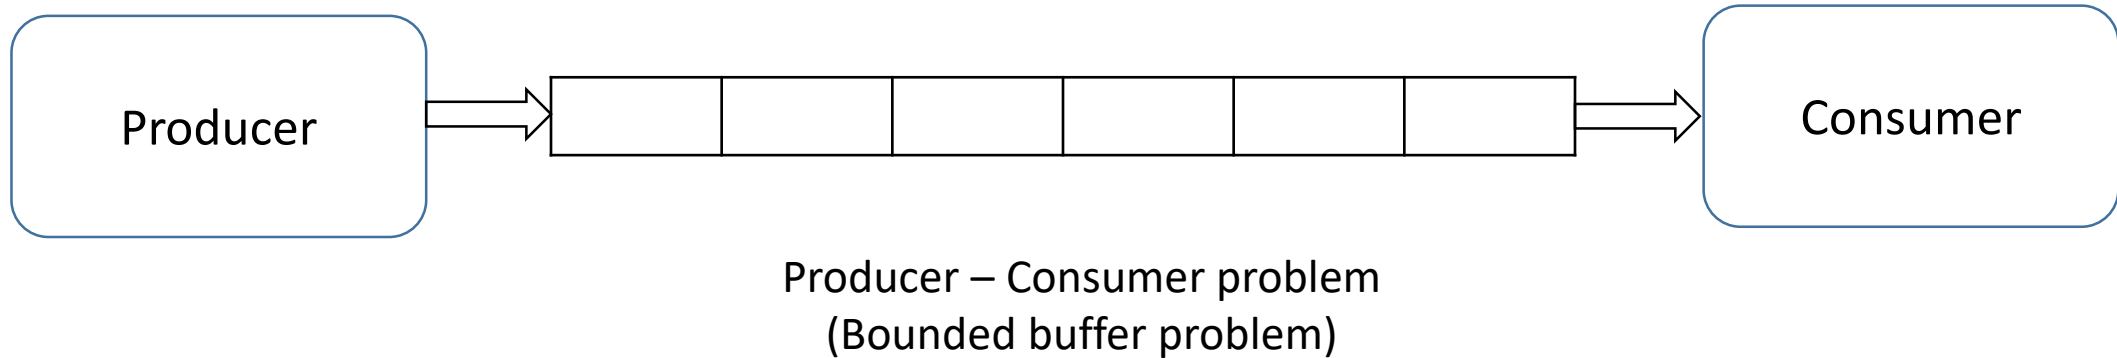
Thundering herd problem

- when a large number of processes (or threads) waiting for an event are woken up when that event occurs
 - but only one process is able to handle the event
- When the processes wake up, they will each try to handle the event, but only one will win
- All waiting processes wakeup simultaneously
- Large number of context switches
- Can lead to starvation

```
void lock(int *locked) {  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
        else  
            add process to queue  
            sleep()  
    }  
}
```

```
void unlock(*locked) {  
    locked = 0;  
    remove process P from queue  
    wakeup(P);  
}
```

Semaphore - another synchronization primitive



- Problem
 - Producer produces fast, but buffer is FULL
 - Consumer consumes fast, but buffer is EMPTY

mutex, full, and empty are mutexes

Producer

```
while(TRUE) {  
    item = produce_item;  
    if (count == N) sleep(empty);  
    lock(mutex);  
    insert_item_to_buffer;  
    count++;  
    unlock(mutex);  
    if(count == 1) wakeup(full);  
}
```

Consumer

```
while(TRUE) {  
    if(count == 0) sleep(full);  
    lock(mutex);  
    item = remove_item_from_buffer;  
    count--;  
    unlock(mutex);  
    if(count == N-1) wakeup(empty);  
    item = consume_item;  
}
```

Problem arises when

Count = 0

Context switch – Producer

```
item = produce_item;  
if (count == N) sleep(empty);  
lock(mutex);  
insert_item_buffer;  
count++;  
unlock(mutex);  
if(count == 1) wakeup(full);
```

Context switch – Consumer

```
if(count == 0) sleep(full);
```


- Consumer uses the old value of count [=0]
 - Consumer waits
- Producer goes on producing
 - Until the buffer is full
- Further, both producer and consumer wait infinitely
- Three mutexes – *full*, *empty*, and *mutex* – are insufficient to solve the problem

Semaphore – proposed by Dijkstra (1965)

- Two atomic functions – down() and up(); also called P() and V()
- S is a shared memory location

```
void down(int *S) {  
    while(*S <= 0);  
    *S--;  
}
```

```
void up(int *S) {  
    *S++;  
}
```



- Two semaphores – empty and full
- Indicate the number of empty and full slots in the buffer respectively

full and empty are semaphores

Producer

```
while(TRUE) {  
    item = produce_item;  
    down(empty);  
    insert_item_to_buffer;  
    up(full)  
}
```

Consumer

```
while(TRUE) {  
    down(full);  
    remove_item_from_buffer;  
    up(empty);  
    item = consume_item;  
}
```

full and empty are semaphores
initially buffer is FULL

Producer

```
while(TRUE) {  
    item = produce_item;  
    down(empty);  
    insert_item_to_buffer;  
    up(full)  
}
```

Consumer

```
while(TRUE) {  
    down(full);  
    remove_item_from_buffer;  
    up(empty);  
    item = consume_item;  
}
```

full and empty are semaphores
initially buffer is EMPTY

Producer

```
while(TRUE) {  
    item = produce_item;  
    down(empty);  
    insert_item_to_buffer;  
    up(full)  
}
```

Consumer

```
while(TRUE) {  
    down(full);  
    remove_item_from_buffer;  
    up(empty);  
    item = consume_item;  
}
```

Synchronized access – using mutex

full and empty are semaphores


Producer

```
while(TRUE) {  
    item = produce_item;  
    down(empty);  
    lock(mutex);  
    insert_item_to_buffer;  
    unlock(mutex);  
    up(full)  
}
```

Consumer

```
while(TRUE) {  
    down(full);  
    lock(mutex);  
    remove_item_from_buffer;  
    unlock(mutex);  
    up(empty);  
    item = consume_item;  
}
```

Dining Philosopher problem



- Philosophers are independent
- At any instant, a philosopher is either eating or thinking
- When a philosopher wants to eat, he uses two forks - one from the left and one from the right
- When a philosopher wants to think, he keeps down both forks at their original place.
- **Problem:** None of the philosophers should STARVE

Possible Solution - 1

```
philosopher(int i) {  
    while(TRUE) {  
        think();  
        pick_fork(R);  
        pick_fork(L);  
        eat();  
        put_fork(L);  
        put_fork(R);  
    }  
}
```

- What happens if only Philosophers P2 and P4 are given priority?
 - P2 uses F2 and F3
 - P4 uses F4 and F5
 - P1, P3, and P5 never get the forks and are starved
- What happens if all philosophers think and eat at the same time with think_time_quantum the same?
 - All philosophers pick their right forks
 - None gets the left fork
 - Deadlock situation: indefinite execution without any progress
- Solution is not ideal

Possible Solution - 2

```
philosopher(int i) {  
    while(TRUE) {  
        think();  
        pick_fork(R);  
        if (available(L)) {  
            pick_fork(L);  
            eat();  
            put_fork(L);  
            put_fork(R);  
        }  
        else {  
            put_fork(R);  
            sleep(T);  
        }  
    }  
}
```

- Situation when all the philosophers execute simultaneously; think and eat at the same time
- Again leads to deadlock as all the philosophers lift the right fork, upon finding the unavailability of the left fork, put back the right fork
- Solution is not ideal
- What if the philosophers sleep for a `random_time` rather than `fixed_time T`?
 - Still starvation cannot be ruled out

Possible Solution – 3 (using mutex)

```
philosopher(int i) {  
    while(TRUE) {  
        think();  
        lock(mutex);  
        pick_fork(R);  
        pick_fork(L);  
        eat();  
        put_fork(L);  
        put_fork(R);  
        unlock(mutex);  
    }  
}
```

- Protection of critical section
- Prevents deadlock
- But the problem/issue is only one philosopher can eat at any given time
 - Since the mutex is shared by all the philosophers

Possible Solution – 3 (using semaphore)

- Uses N semaphores (s[1], s[2], ... s[N]) all initialized to zero
- Philosopher can be in three states; HUNGRY, EATING, and THINKING
- A philosopher can EAT only if the neighbours do not eat

```
void philosopher(int i) {  
    while(TRUE) {  
        think();  
        pick_forks(i);  
        eat();  
        put_forks(i);  
    }  
}
```

```
void pick_forks(int i) {  
    lock(mutex);  
    state[i] = HUNGRY;  
    test(i);  
    unlock(mutex);  
    down(s[i]);  
}
```

```
void put_forks(int i) {  
    lock(mutex);  
    state[i] = THINKING;  
    test(LEFT);  
    test(RIGHT);  
    unlock(mutex);  
}
```

```
void test(int i) {  
    if(state[i] = HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {  
        state[i] = EATING;  
        up(s[i]);  
    }  
}
```

	P1	P2	P3	P4	P5
State	T	T	T	T	T
Semaphore	0	0	0	0	0

	P1	P2	P3	P4	P5
State	T	H	T	T	T
Semaphore	0	1	0	0	0

	P1	P2	P3	P4	P5
State	T	E	T	T	T
Semaphore	0	0	0	0	0

	P1	P2	P3	P4	P5
State	T	E	H	T	T
Semaphore	0	0	0	0	0

P3 gets blocked

	P1	P2	P3	P4	P5
State	T	T	H	T	T
Semaphore	0	0	0	0	0

	P1	P2	P3	P4	P5
State	T	T	E	T	T
Semaphore	0	0	1	0	0

	P1	P2	P3	P4	P5
State	T	T	E	T	T
Semaphore	0	0	0	0	0


Deadlocks

- System consists of resources
- Resource types R_1, R_2, \dots, R_m
CPU cycles, memory space, I/O devices
- Each resource type R_i has W_i instances.
- Each process utilizes a resource as follows:
 - request
 - use
 - release


Ref.: https://www.cs.uic.edu/~jbell/CourseNotes/OperatingSystems/7_Deadlocks.html

Resource-Allocation Graph


- A set of vertices V and a set of edges E .
- V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system
- **request edge** – directed edge $P_i \rightarrow R_j$
- **assignment edge** – directed edge $R_j \rightarrow P_i$




- Process





- Resource Type with 4 instances




- P_i requests instance of R_j



- P_i is holding an instance of R_j




Resource Allocation Graph With a Deadlock



- Presence of a cycle in a graph is a potential deadlock situation
- Need not always end in a deadlock situation, but probably can end up

Graph With A Cycle But No Deadlock



- If graph contains no cycles \Rightarrow no deadlock
- If graph contains a cycle \Rightarrow
 - if only one instance per resource type, then deadlock
 - if several instances per resource type, possibility of deadlock

Deadlock Characterization

- Deadlock can arise if the following four conditions hold simultaneously:
 - **Mutual exclusion:** only one process at a time can use a resource
 - **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes
 - **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task
 - **Circular wait:** there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .

Handling deadlocks

- Ensure that the system will never enter a deadlock state:
 - Deadlock prevention
 - Deadlock avoidance
- Allow the system to enter a deadlock state and then recover

Deadlock Prevention

- **Mutual Exclusion** – not required for sharable resources (e.g., read-only files); must hold for non-sharable resources
- **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
 - Require process to request and be allocated all its resources before it begins execution, or allow process to request resources only when the process has none allocated to it.
 - Low resource utilization; starvation possible

Deadlock Prevention ... Contd.

- **No Preemption** –
 - If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released
 - Preempted resources are added to the list of resources for which the process is waiting
- **Circular Wait** – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration

Deadlock Avoidance


- Requires that the system has some additional a priori information available
 - Simplest and most useful model requires that each process declare the maximum number of resources of each type that it may need
 - The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition
 - Resource-allocation state is defined by the number of available and allocated resources, and the maximum demands of the processes

Safe State

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state
- System is in safe state if there exists a sequence $\langle P_1, P_2, \dots, P_n \rangle$ of ALL the processes in the systems such that for each P_i , the resources that P_i can still request can be satisfied by currently available resources + resources held by all the P_j , with $j < i$
- That is:
 - If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished
 - When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate
 - When P_i terminates, P_{i+1} can obtain its needed resources, and so on

Safe State ... Contd.

- If a system is in safe state \Rightarrow no deadlocks
- If a system is in unsafe state \Rightarrow possibility of deadlock
- Avoidance \Rightarrow ensure that a system will never enter an unsafe state.




Deadlock Avoidance Algorithms

- Single instance of a resource type
 - Use a resource-allocation graph
- Multiple instances of a resource type
 - Use the Banker's algorithm


Resource-Allocation Graph Scheme

- **Claim edge** $P_i \rightarrow R_j$ indicated that process P_j may request resource R_j ; represented by a dashed line
- **Claim edge converts to request edge** when a process requests a resource
- **Request edge converted to an assignment edge** when the resource is allocated to the process
- When a resource is released by a process, assignment edge reconverts to a claim edge
- Resources must be claimed *a priori* in the system

Claim edge



Unsafe state



Suppose that process P_i requests a resource R_j

The request can be granted only if **converting the request edge to an assignment edge does not result in the formation of a cycle** in the resource allocation graph

Banker's Algorithm

- Multiple instances
- Each process must indicate a priori the maximum resource usage
- When a process requests a resource, it may have to wait
- When a process gets all its resources it must return them in a finite amount of time

Banker's Algorithm ... Contd.

Let n = number of processes, and m = number of resources types

- **Available:** Vector of length m . If available $[j] = k$, there are k instances of resource type R_j available
- **Max:** $n \times m$ matrix. If $Max[i,j] = k$, then process P_i may request at most k instances of resource type R_j
- **Allocation:** $n \times m$ matrix. If $Allocation[i,j] = k$ then P_i is currently allocated k instances of R_j
- **Need:** $n \times m$ matrix. If $Need[i,j] = k$, then P_i may need k more instances of R_j to complete its task

$$Need[i,j] = Max[i,j] - Allocation[i,j]$$

Safety Algorithm

Let Work and Finish be vectors of length m and n, respectively. Initialize:

Work = Available

Finish [i] = false for $i = 0, 1, \dots, n-1$

1. Find an i such that both:

(a) Finish [i] = false

(b) $\text{Need}_i \leq \text{Work}$

If no such i exists, go to step 4

2. $\text{Work} = \text{Work} + \text{Allocation}_i$

Finish[i] = true

go to step 2

If Finish [i] == true for all i, then the system is in a safe state

Resource-Request Algorithm for Process P_i

$Request_i$ = request vector for process P_i . If $Request_i[j] = k$ then process P_i wants k instances of resource type R_j

1. If $Request_i \leq Need_i$ go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
2. If $Request_i \leq Available$, go to step 3. Otherwise P_i must wait, since resources are not available
3. Pretend to allocate requested resources to P_i by modifying the state as follows:

$$\begin{aligned} Available &= Available - Request_i \\ Allocation_i &= Allocation_i + Request_i \\ Need_i &= Need_i - Request_i \end{aligned}$$

- If safe \Rightarrow the resources are allocated to P_i
- If unsafe $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored

Example

- 5 processes P_0 through P_4 ; 3 resource types: A (10 instances), B (5 instances), and C (7 instances)
- Snapshot at time T_0 :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>	<u>Need</u> (Max-Allocation)
	A B C	A B C	A B C	A B C
P_0	0 1 0	7 5 3	3 3 2	7 4 3
P_1	2 0 0	3 2 2		1 2 2
P_2	3 0 2	9 0 2		6 0 0
P_3	2 1 1	2 2 2		0 1 1
P_4	0 0 2	4 3 3		4 3 1

Is the sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ in a SAFE STATE?

Work = Avail = [3 3 2]

Finish[5] = {F, F, F, F, F}

P1

Need[1] <= Work & Finish[1] = F

[1 2 2] <= [3 3 2]

Work = Work + Alloc[1]

Work = [3 3 2] + [2 0 0] = [5 3 2]

Finish[1] = T;

P2

Need[2] <= Work & Finish[2] = F

[6 0 0] <= [7 4 5]

Work = Work + Alloc[1]

Work = [7 4 5] + [3 0 2] = [10 4 7]

Finish[2] = T;

P3

Need[3] <= Work & Finish[3] = F

[0 1 1] <= [5 3 2]

Work = Work + Alloc[3]

Work = [5 3 2] + [2 1 1] = [7 4 3]

Finish[3] = T

P0

Need[0] <= Work & Finish[0] = F

[7 4 3] <= [10 4 7]

Work = Work + Alloc[0]

Work = [7 4 3] + [2 1 1] = [9 5 4]

Finish[0] = T

P4

Need[4] <= Work & Finish[4] = F

[4 3 1] <= [7 4 3]

Work = Work + Alloc[3]

Work = [7 4 3] + [0 0 2] = [7 4 5]

Finish[4] = T

The sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ is in a SAFE STATE

Can request for (1,0,2) by P1 be granted?

$\text{req}[1] \leq \text{need}[1]$

$\text{req}[1] \leq \text{avail}$

$[1\ 0\ 2] \leq [3\ 3\ 2]$

$[1\ 0\ 2] \leq [3\ 3\ 2]$

$\text{avail} = \text{avail} - \text{req}[1]$

$\text{alloc}[1] = \text{alloc}[1] + \text{req}[1]$

$\text{need}[1] = \text{need}[1] - \text{req}[1]$

$\text{avail} = [3\ 3\ 2] - [1\ 0\ 2] = [2\ 3\ 0]$

$\text{alloc}[1] = [2\ 0\ 0] + [1\ 0\ 2] = [3\ 0\ 2]$

$\text{need}[1] = [1\ 2\ 2] - [1\ 0\ 2] = [0\ 2\ 0]$

Safe state; request granted



	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P ₀	0 1 0	7 4 3	2 3 0
P ₁	3 0 2	0 2 0	
P ₂	3 0 2	6 0 0	
P ₃	2 1 1	0 1 1	
P ₄	0 0 2	4 3 1	

- Executing safety algorithm does the sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfy safety requirement?
- Can request for (3,3,0) by **P₄** be granted?
- Can request for (0,2,0) by **P₀** be granted?

Deadlock Detection

- Single Instance of Each Resource Type
 - Maintain **wait-for** graph
 - Nodes are processes
 - $P_i \rightarrow P_j$ if P_i is waiting for P_j
 - Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock
 - An algorithm to detect a cycle in a graph requires an order of n^2 operations, where n is the number of vertices in the graph

Resource-Allocation Graph and Wait-for Graph



Detection algorithm

1. Let Work and Finish be vectors of length m and n , respectively
 - (a) Work = Available
 - (b) For $i = 1, 2, \dots, n$, if $\text{Allocation}_i \neq 0$, then
Finish[i] = false; otherwise, Finish[i] = true
2. Find an index i such that both:
 - (a) Finish[i] == false
 - (b) $\text{Request}_i \leq \text{Work}$If no such i exists, go to step 4
3. $\text{Work} = \text{Work} + \text{Allocation}_i$ and set Finish[i] = true
Go to step 2
4. If Finish[i] == false, for some i , $1 \leq i \leq n$, then the system is in deadlock state. Moreover, if Finish[i] == false, then P_i is deadlocked

Example

- Five processes P_0 through P_4 ; three resource types; A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time T_0 :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	0 0 0	0 0 0
P_1	2 0 0	2 0 2	
P_2	3 0 3	0 0 0	
P_3	2 1 1	1 0 0	
P_4	0 0 2	0 0 2	

- Sequence $\langle P_0, P_2, P_3, P_1, P_4 \rangle$ will result in $\text{Finish}[i] = \text{True}$ for all i

Sequence $\langle P_0, P_2, P_3, P_1, P_4 \rangle$

Work = avail = [0 0 0]
Finish[5] = {F, F, F, F, F}

P0

req[0] ≤ work
work = work + alloc[0]
Finish[5] = {T, F, F, F, F}

Y; [0 0 0] ≤ [0 0 0]
work = [0 0 0] + [0 1 0] = [0 1 0]

P1

req[1] ≤ work
work = work + alloc[1]
Finish[5] = {T, T, T, T, F}

Y; [2 0 2] ≤ [5 2 4]
work = [5 2 4] + [2 0 0] = [7 2 4]

P2

req[2] ≤ work
work = work + alloc[2]
Finish[5] = {T, F, T, F, F}

Y; [0 0 0] ≤ [0 1 0]
work = [0 1 0] + [3 0 3] = [3 1 3]

P4

req[4] ≤ work
work = work + alloc[4]
Finish[5] = {T, T, T, T, T}

Y; [0 0 2] ≤ [7 2 4]
work = [7 2 4] + [0 0 2] = [7 2 6]

P3

req[3] ≤ work
work = work + alloc[3]
Finish[5] = {T, F, T, T, F}

Y; [1 0 0] ≤ [3 1 3]
work = [3 1 3] + [2 1 1] = [5 2 4]

Deadlock Recovery

- Manual intervention
 - Abort all deadlocked processes
- Terminate one or more processes involved in the deadlock
- Preempt resources
 - Selecting a victim process
 - Rollback – return to some safe state, restart process for that state
 - Starvation – same process may always be picked as victim

Completely Fair Scheduler (CFS)

- It is based on Rotating Staircase Deadline Scheduler (RSDL).
- It is default scheduling process since version 2.6.23.
- Elegant handling of I/O and CPU bound process.
- It fairly or equally divides the CPU time among all the processes
- If there are N processes in the ready queue then each process receives $(100/N)\%$ of CPU time according to CFS.

Example

Process	Burst Time (ms)
A	10
B	6
C	14
D	6

A = 4
C = 8

A	1	2	3	4	5	6	8	10	
B	1	2	3	4	5	6			
C	1	2	3	4	5	6	8	10	14
D	1	2	3	4	5	6			

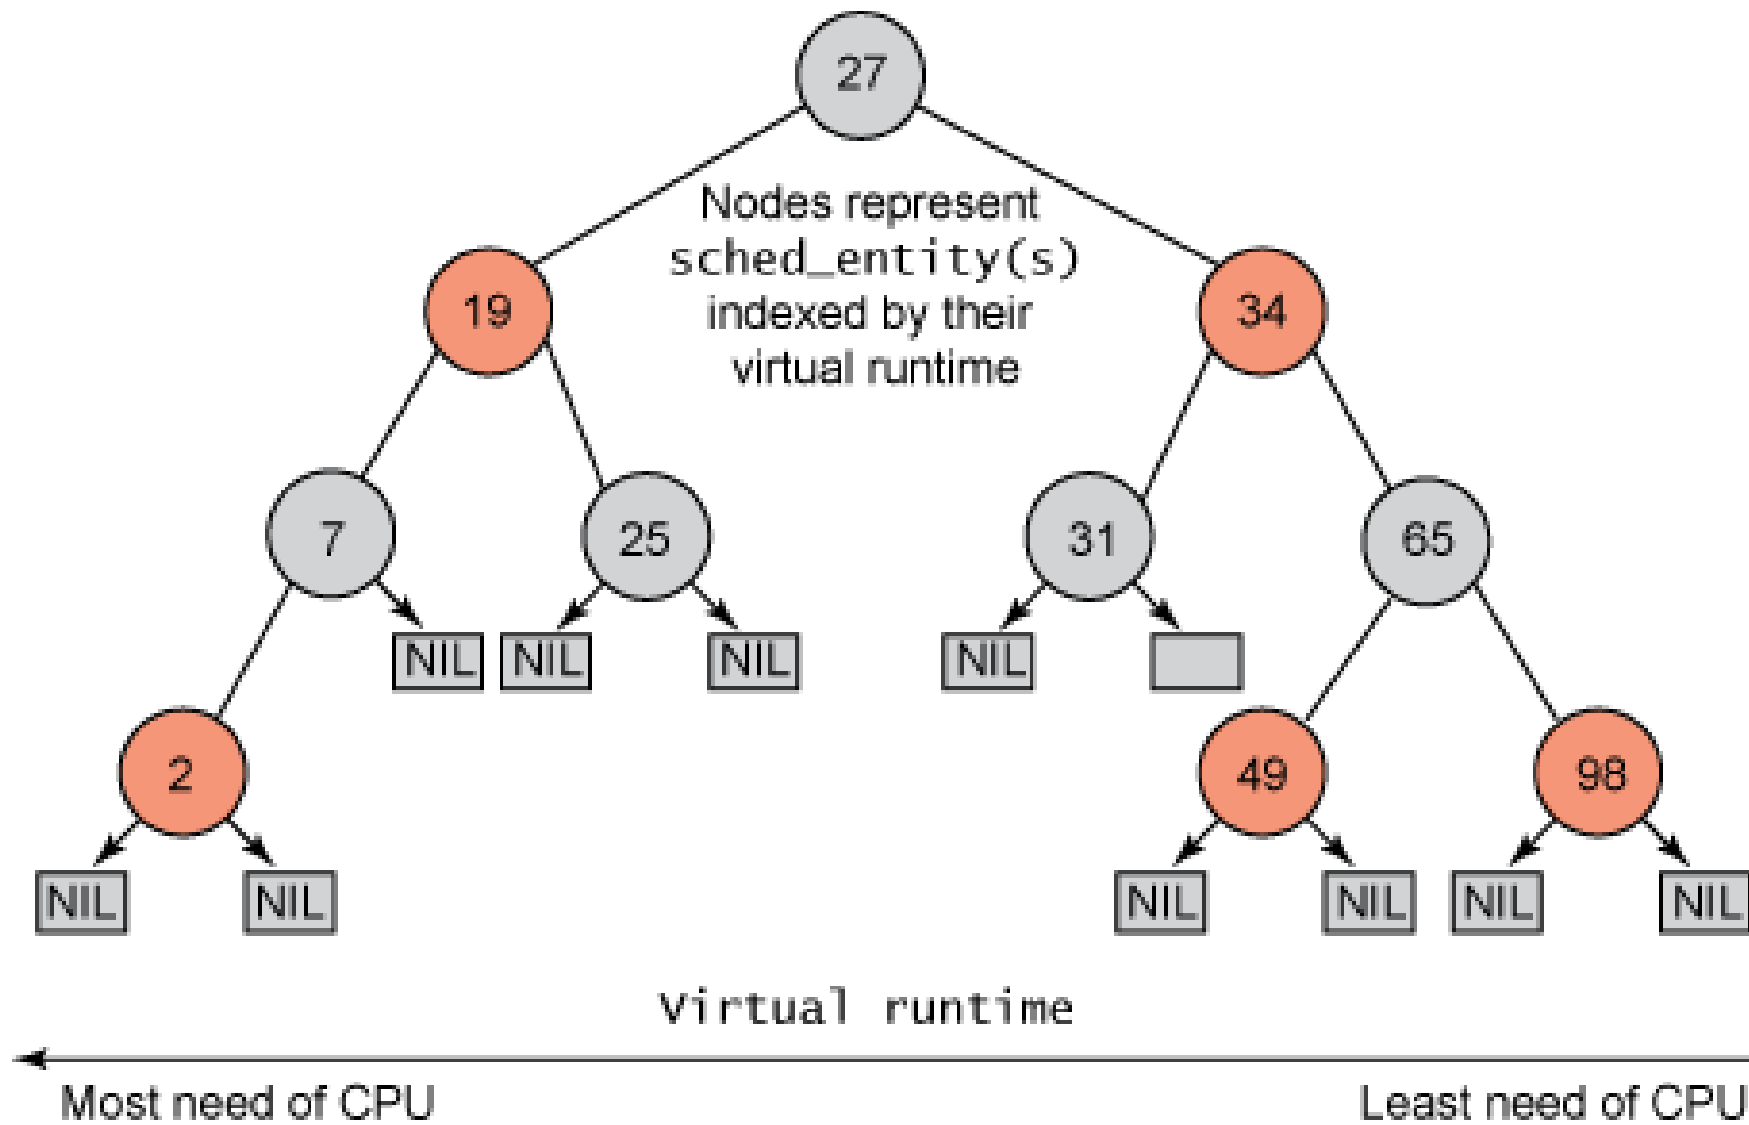
C = 4

- Let time quanta be 4 ms
- Each process gets, $4/4 = 1$ ms
- After 6 ms, only two processes left; each process gets $4/2 = 2$ ms
- After 10 ms, only process C is left; Process C gets the entire time quanta of 4 ms.

Virtual Runtime

- Each runnable process has a virtual time associated with it in Process Control Block (PCB)
- Whenever a context switch happens (or at every scheduling point) then current running process virtual time is increased by $\text{virtualruntime_currprocess} += T$; where T is time for which it is executed recently.
- Runtime for the process therefore monotonically increases.
- During interrupt, process with the lowest virtual runtime is chosen

implemented using RED-BLACK Trees and not Queues



- When the scheduler is invoked to run a new process:
 - The leftmost node of the scheduling tree is chosen and sent for execution
 - If the process simply completes execution, it is removed from the system and scheduling tree
 - If the process reaches its *maximum execution time** or is otherwise stopped (voluntarily or via interrupt) it is reinserted into the scheduling tree based on its new spent *execution time*
 - The new leftmost node will then be selected from the tree, repeating the iteration

* *the time the process would have expected to run on an "ideal processor"*

- all the process which are in main memory are inserted into Red Black trees and whenever a new process comes it is inserted into the tree
- During context switch
 - The virtual time for the current process which was executing is updated
 - The new process is decided which has lowest virtual time and it is the left most node of Red Black tree
 - If the current process still has some burst time then it is inserted into the Red Black tree

wait_runtime

- As a process waits for the CPU, the scheduler tracks the amount of time it would have used on the ideal processor
- is used to rank processes for scheduling and to determine the amount of time the process is allowed to execute before being preempted
- The process with the longest wait time is picked by the scheduler and assigned to the CPU
- Process running – wait time decreases

CFS and Priority

- implements priorities by using weighted tasks
 - each task is assigned a weight based on its static priority
- $\text{vruntime} += t * (\text{weight of a process})$
- the task with lower weight (lower-priority) will see time elapse at a faster rate than that of a higher-priority task
- `wait_runtime` will exhaust more quickly than that of a higher-priority task, so lower-priority tasks will get less CPU time compared to higher-priority tasks

- I/O bound processes (interactive processes) should get higher priority than CPU bound processes (batch processes)
- I/O bound processes -> small CPU bursts -> low vruntime -> appears on the left of RBTree -> higher priority